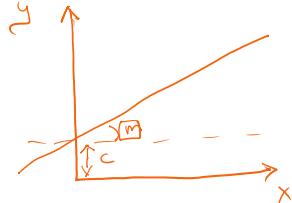


Gen AI - Generate Content

AI - Artificial Intelligence

ML - Machine learning → understand data

DL - Deep learning → processing power



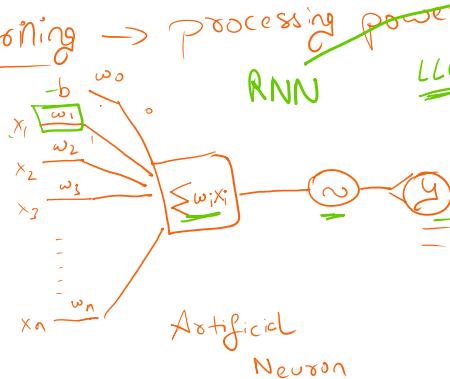
$$y = mx + c$$

Linear deg.

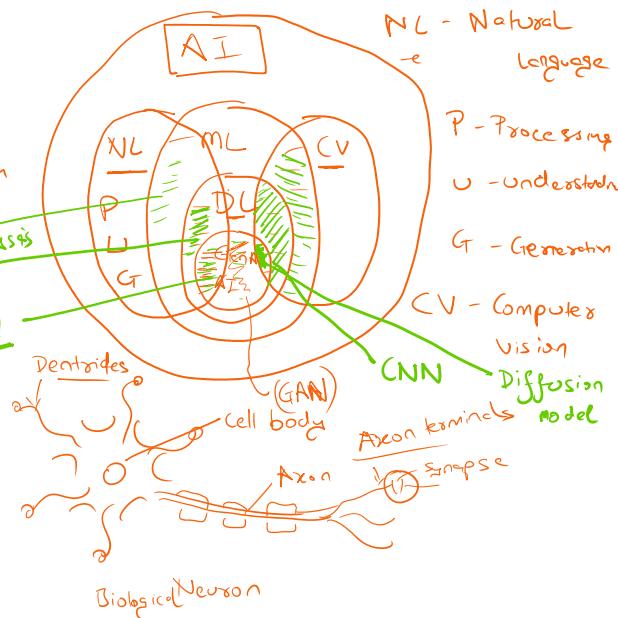
$$y = m_1x_1 + m_2z_2 + \dots + m_nx_n + c$$

$$y = m_1x_1 + m_2z_2$$

Data collect.,

 ω = weights b = bias, c = constan., w_0 x_i = feature y = output

Ⓐ - Activation function

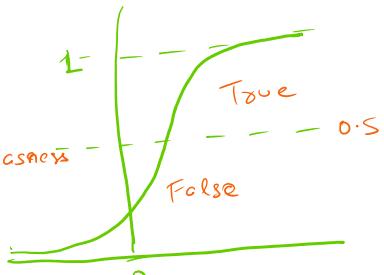
Activation function

- Step func.

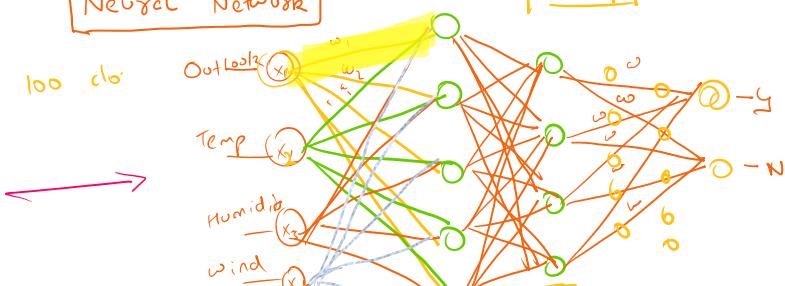
$$\boxed{\text{Sigmoid func}} = \frac{1}{1 + e^{-z}} = [0, 1]$$

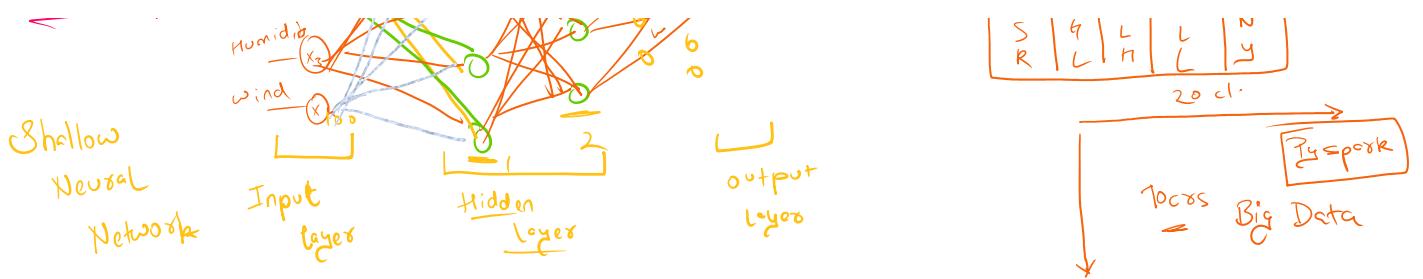
- ReLU

- tanh

 x = feature ω = weights b = (helper value) biasnessBuilding block
of DL

Artificial Neuron
Precptone

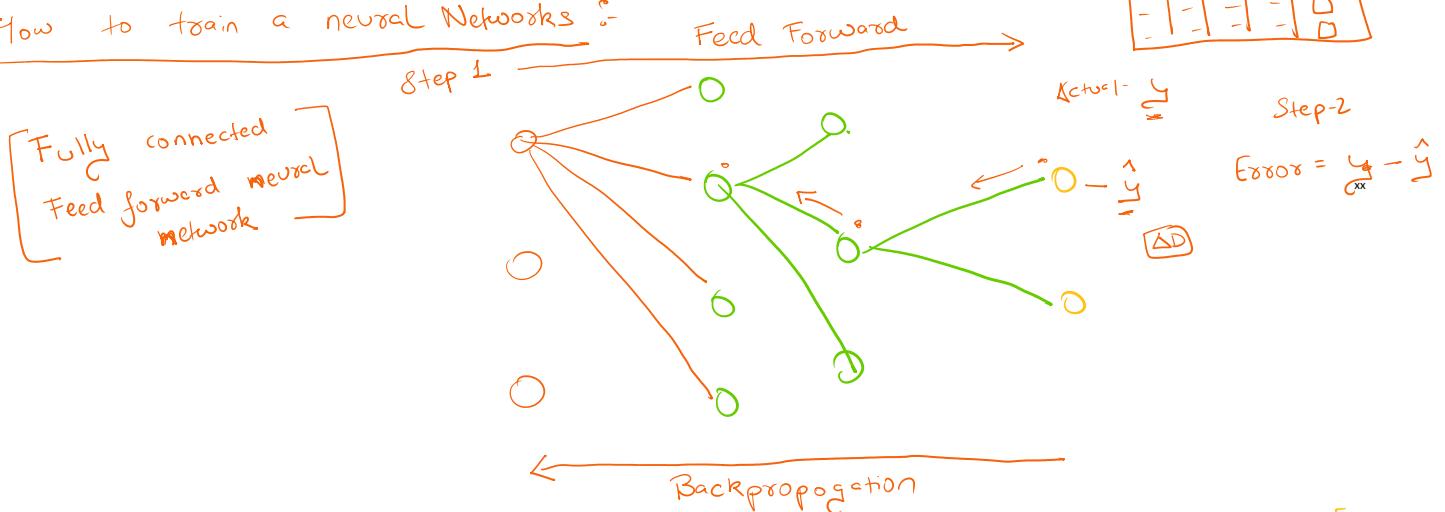
Neural NetworkDeep Neural Network



Why we need Deep Learning instead Machine learning ?

- Machine
 - ◦ Structured data
 - ◦ Feature Engg.
 - Not good for high dimensional data
- Deep learning
 - can work unstructured data - audio, text, img. - High Dimensional data
 - High dimensional data
 - Automatic Feature Selection
 - High computation cost

How to train a neural Network :-



Feed Forward

$$\begin{aligned} z &= \sum x_i w_i \\ y &= f(z) = \frac{1}{1+e^{-z}} \end{aligned}$$

$$\begin{aligned} \text{Step-1} \\ a_1 &= (\omega_{13} \times x_1) + (\omega_{23} \times x_2) \\ &= (0.1 \times 0.35) + (0.8 \times 0.8) = 0.755 \end{aligned}$$

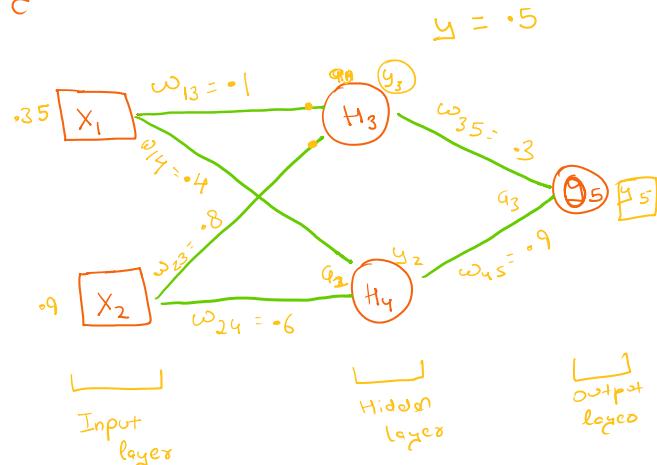
$$y_3 = \frac{1}{1+e^{-0.755}} = 0.68$$

$$a_2 = (0.4 \times 0.35) + (0.9 \times 0.6) = 0.68$$

$$y_4 = \frac{1}{1+e^{-0.68}} = 0.6637$$

$$a_3 = (0.68 \times 0.3) + (0.6637 \times 0.9) = 0.8013$$

$$y_5 = \frac{1}{1+e^{-0.8013}} = 0.6902$$



Step-2

$$\begin{aligned} \text{Error} &= \hat{y} - y \\ 0.5 - 0.69 &= -0.19 \end{aligned}$$

Step -3

Backpropagation

$$\Delta \omega_{ji} = \eta * \delta_j O_i$$

$$\delta_j = O_j(1-O_j)(t_j - O_j) \quad \text{if } j \text{ is output}$$

$$\delta_j = O_j(1-O_j) \geq \delta_k \omega_{kj} \quad \text{if } j \text{ is hidden neuron}$$

η = learning rate

δ_j = error measure of joint

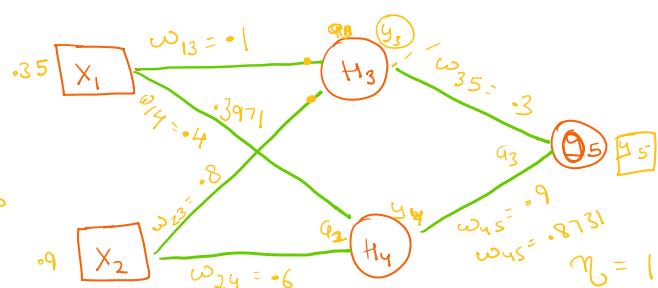
t_j = correct output

O = output of neuron

For output neuron

$$\delta_s = \hat{y}(1-\hat{y})(y - \hat{y})$$

$$= 0.6902(1-0.6902)(0.5 - 0.6902) = -0.0406$$



$$\delta_3 = y_3(1-y_3)(\omega_{35} \times \delta_s)$$

$$= 0.68(1-0.68)(0.3 \times -0.0406) = -0.00265$$

$$\delta_4 = y_4(1-y_4)(\omega_{45} \times \delta_s) =$$

$$0.6637(1-0.6637)(0.9 \times -0.0406) = -0.00815$$

Updating weights

$$\Delta \omega_{di} = \eta \delta_i y_i$$

$$\Delta \omega_{45} = 1 \times -0.0406 \times 0.6637 = -0.0269$$

$$(\text{new}) \omega_{45}^{(1)} = \Delta \omega_{45} + \omega_{45} (\text{old}) = -0.0269 + 0.9 = 0.8731$$

$$\Delta \omega_{14} = \eta \delta_4 y_4$$

$$= 1 \times -0.0082 \times 0.35 = -0.0287$$

$$\omega_{14} = 0.3971$$

$$\begin{aligned} \Delta \omega_{45} &= -0.0269 \\ \rightarrow \Delta \omega_{45} &= -0.0026 \\ \rightarrow \Delta \omega_{45} &= -0.000034 \end{aligned}$$

-2nd Iteration

-0.19

→ Feed Forward

→ Error - 0.15

→ Backpropagation

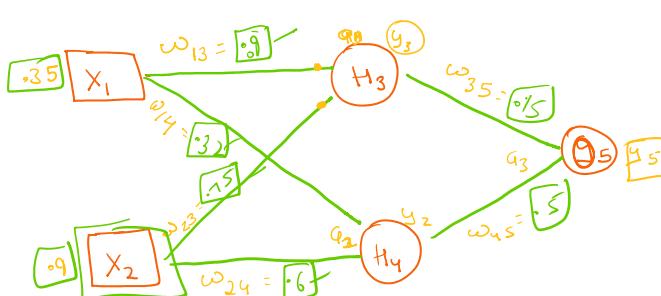
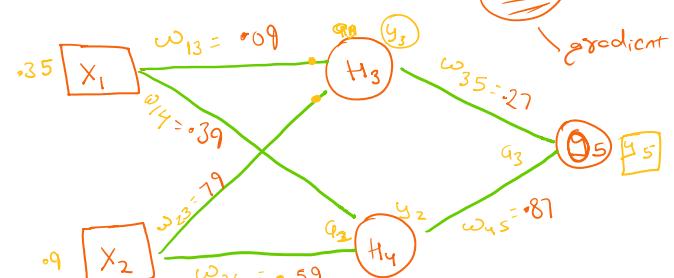
3rd Iteration

Feed Forward

$$\rightarrow E_{\text{Error}} = 0.09 \rightarrow 14.9 \downarrow$$

Y_{out}

$$E_{\text{Error}} = 0.1495$$



Different types of Neural Networks:

- o
- c

Type of Neural Networks

- Fully Connected Neural Network :- (FCNN)
 - ◦ Convolutional Neural Network :- (CNN)
 - ◦ Recurrent Neural Network (RNN)
 - ◦ Long Short term memory : (LSTM)
- ★ Interview question
- GAN - Generative Adversarial Network
 - Autoencoder
- 37 connection

→ Different Architecture of Neural Network :-

◦ one input one output :-

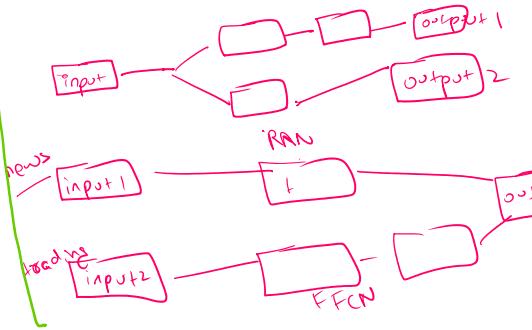


◦

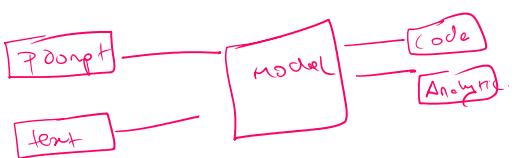
◦ One input multiple output :-



◦ multiple input one output :-



◦ multiple input multiple output :-



$$\Delta w = -0.026$$

$$\Delta w = 1$$

→ Activation function :-

◦ Sigmoid :- $\frac{1}{1 + e^{-z}}$

Range :- (0 to 1)

→ Vanishing Gradient

→ Exploding Gradient

→ use case:-

- binary classification

→ Pros :-

probability based values which are interpretable

→ Cons:-

- Vanishing Gradient

◦ Hyperbolic tan → $\tanh = \tanh(x) = \frac{e^x - e^{-x}}{e^x + e^{-x}}$

Range : $[-1, 1]$

Use Cases:-

- Generally used in hidden layer

Pros:-

- Zero centered output helps optimizer
- less likely to saturate

Cons:-

Vanishing gradient

→ ReLU

$$\text{ReLU}(x) = \max(0, x)$$

Rectified Linear Unit

Range = $[0, \infty)$

Use Cases :-

Default activation for hidden layers

Pros:-

- Computationally efficient
- Mitigates vanishing gradient

Cons:-

Dying ReLU Problem , Not zero centered

4) Leaky ReLU $[0, \infty)$

Parametric ReLU pReLU

ELU - Exponential Linear Unit

$$LR = \begin{cases} x & \text{if } x > 0 \\ \alpha x & \text{if } x \leq 0 \end{cases}$$

$$\begin{cases} x & \text{if } x > 0 \\ \alpha x & \text{if } x \leq 0 \end{cases}$$

$\alpha =$ is trained

$$\begin{cases} x & \text{if } x > 0 \\ \alpha(e^x - 1) & \text{if } x \leq 0 \end{cases}$$

$$\text{LR} - \begin{cases} x & \text{if } x \leq 0 \\ 0 & \text{if } x > 0 \end{cases} \rightarrow \begin{cases} x & \text{if } x \leq 0 \\ 0 & \text{if } x > 0 \end{cases} \rightarrow \infty \text{ if } x = \infty$$

$x =$ very small number $x =$ is defined

Range = $(-\infty, \infty)$

Use Case's

to solve Dying ReLU Problem.

Posse

Prevent dead neurons

Const not zero centered
 α is tricky

α is Adoptive

Cons

increase Model complexity

in deep networks
to improve learning
processes

Pros
Dying Relu
Vanishing

Complex

- swish :-

$$= \frac{x}{1+e^{-x}} = x \circ \sigma(x) - 10 \times (0, 1) - vc$$

$x \neq 0 \Rightarrow 0$
 $-1000000000 \times 000001$

$$\text{Range} : (-\infty, \infty)$$

User

deep network

Poos !
- Avoid zero gradient

Com. Computable ens

$$\text{softmax} \stackrel{?}{=} \frac{e^{x_i}}{\sum e^{x_j}}$$

$$\text{Range} = [0, 1]$$

use cos

~~Y~~: multiclass classification

Pros

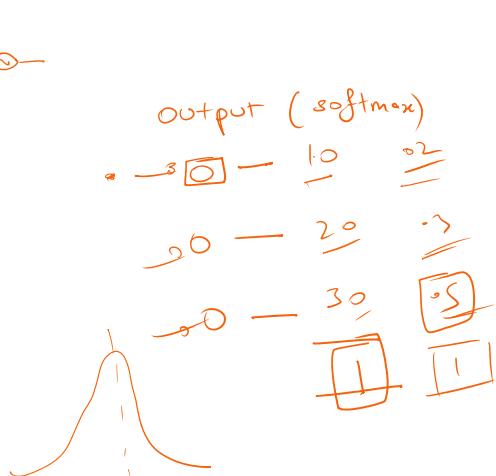
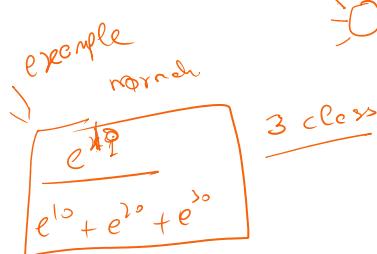
Convert logits to prob. distribution

Cong

- outliers date -

- called them
= class are mut

- class re mutation exclusive -



Cons - outliers don't -
- class are mutually exclusive -

→ Hard Sigmoid

$$= \min_{\bar{x}} \left(1, \max_{x} (0, 0.2x + 0.5) \right)$$

$$\text{Range} = (0, 1)$$

Use:

faster approach value

Project

computationally faster

Const

less smooth than regular sigmoid

→ soft plus

$$= \log(1 + e^x)$$

Range $\rightarrow (0, \infty)$

User

smooth approx of Relu

Poos

— No dead neurons

Coas 81200

- optimizer

- Gradient Descent is (GD)

$$\theta_{\text{new}} = \theta_{\text{old}} - \eta \nabla J(\theta)$$

- o Use

small datasets

• Pros :-

Simple & effective on convex problem

o Const:

\vdash slow convergences & computationally expensive

θ = weight
 η = learning rate
 $\nabla \delta(\theta)$ = loss function



- Cons: slow convergences & computationally expensive on large dataset

- Stochastic GD : (SGD)

$$\theta = \theta - \eta \nabla f(\theta; x_i, y_i)$$

(x_i, y_i) - individual sample

Use Case :-

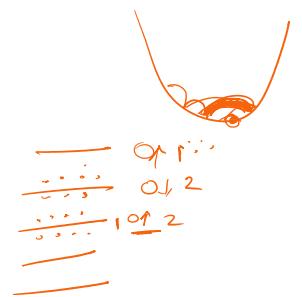
Suitable for large datasets

Pros:

Fast processing & efficient on large dataset

Cons:

Noisy update, may not reach global minima.



- Mini-batch GD :

batch - (mini-batch)

$$\theta = \theta - \eta \nabla f(\theta; \text{batch})$$

Use:

hybrid SGD & GD for large dataset, more stable than SGD

Pros:

- Reduce Noise & variance SGD
- Faster than GD

Cons:

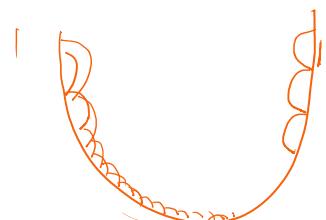
Very large dataset

→ Momentum

$$\vec{v} = \gamma \vec{v} + \eta \nabla f(\theta)$$

$$\theta = \theta - \vec{v}$$

γ = momentum factor =



Use:

useful in training deep network

Pros:

Speeds up the convergences

Pros:

- Speeds the convergence
- Reducing the oscillation

Const:

- Overshoot if momentum is high
- Requires tuning of γ

→ NAG - Nesterov Accelerated Gradient

$$\underline{v} = \underline{v} + \eta \nabla J(\theta - \underline{v})$$

$$\theta = \theta - \underline{v}$$

Use:

it looks ahead before computing the gradient

Pros:

- better than momentum
- prevent overshooting

Const:

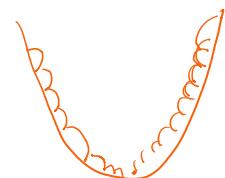
- more complex.

→ Adagrad

$$\theta = \theta - \frac{\eta}{\sqrt{G_t + \epsilon}} \cdot \nabla J(\theta)$$

G_t = sum of sq of past gradient

ϵ = small constant to avoid zero division.



User:

Suitable for sparse data like NLP, image

Pros:

Adapting learning rate based on freq. of parameter



Const:

Learning rate keeps decaying.

- RMSprop : Root Mean Square propagation

$$G_t = \beta G_t + (1 - \beta)(\nabla J(\theta))^2$$

$$\theta = \theta - \frac{\eta}{\sqrt{G_t + \epsilon}} \cdot \nabla J(\theta)$$

$$\beta = \text{decay rate} = 0.9$$

$$\begin{array}{c} \text{---} \\ \text{---} \\ \text{---} \end{array} \quad \begin{array}{c} \uparrow 0.9 \\ \uparrow 0.9 \\ 0.9 \end{array}$$

Pros:

Addressing the adagrad decay prob.

Use :-

Good for non-stationary prob. like in RNN's

Cons :-

Sensitive to β value

- Adam :- Adaptive moment Estimation

First moment estimation

$$m_t = \beta_1 m_t + (1 - \beta_1) \nabla J(\theta)$$

$$\beta_1 = 0.9 - 0.99$$

Second moment estimation

$$v_t = \beta_2 v_t + (1 - \beta_2) \nabla^2 J(\theta)$$

Bias correction :-

$$\hat{m}_t = \frac{m_t}{1 - \beta_1}, \quad \hat{v}_t = \frac{v_t}{1 - \beta_2}$$

Update :-

$$\theta = \theta - \frac{\eta}{\sqrt{\hat{v}_t}} \cdot \hat{m}_t$$

Use case :-

widely used in almost scenarios.

Pros :-

- Combining adaptive learning & momentum learning
- Efficiently works with large datasets & noisy data.
- Faster convergence

Cons :-

- Multiple tunable parameters
- it can overshoot sometimes

Adamax :-

$$v_t = \max(\beta_2 v_{t-1}, |\nabla J(\theta)|)$$

$$\theta = \theta - \frac{\eta}{v_t} \cdot m_t$$

Use :-

very large dataset

Pros :-

more stable updates

Cons

Nadam - Nesterov-accelerated adam

$$m_t = \beta_1 m_t + (1 - \beta_1) \nabla J(\theta)$$

$$\theta = \theta - \frac{\eta}{\sqrt{v_t} + \epsilon} \cdot m_t$$

Use + Pros

adaptive learning of adam + fast convergence nature of NAG

Cons

More sensitive to tuning η parameters

Adamax

$$m_t = \beta m_{t-1} + (1-\beta) \nabla J(\theta)$$

$$v_t = \max(\beta_2 v_{t-1}, |\nabla J(\theta)|)$$

$$= \theta - \frac{\eta}{v_t} \cdot m_t$$

Use +

B) Better for model that sparse updates

Pros

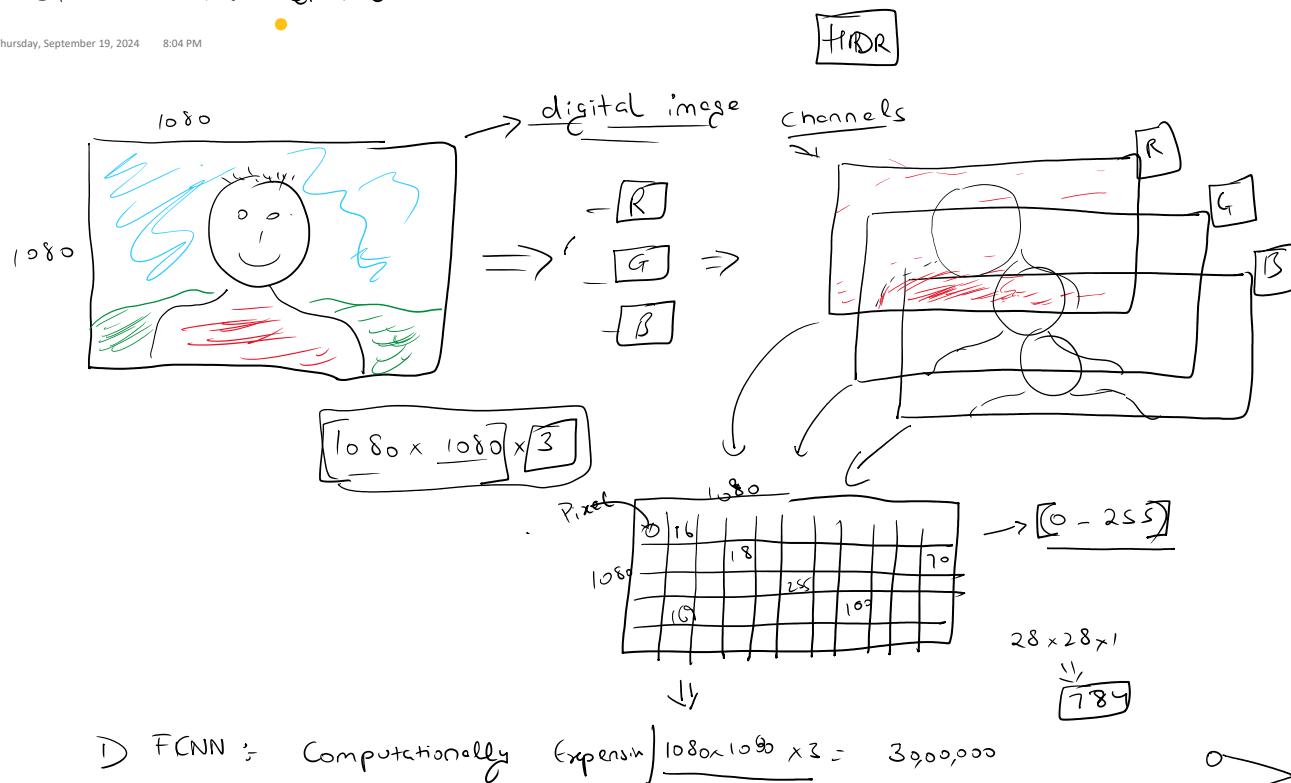
more stable than Adam

Cons

slightly slower than Adam.

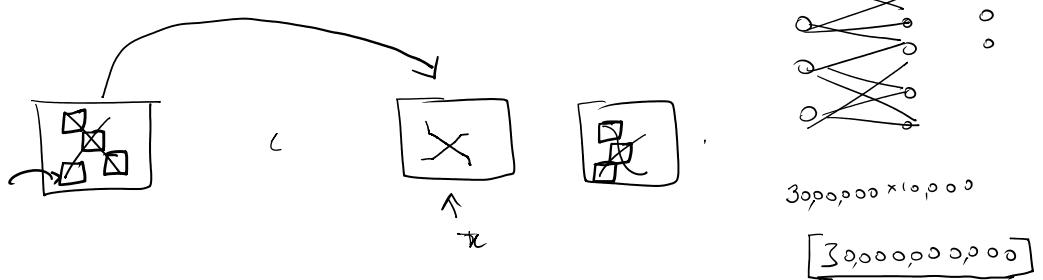
CNN - Convolutional Neural Networks

Thursday, September 19, 2024 8:04 PM



▷ F CNN : Computationally expensive $|1080 \times 1080 \times 3| = 3,900,000$

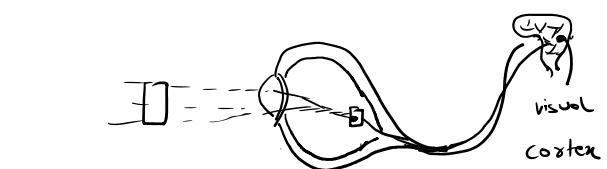
▷

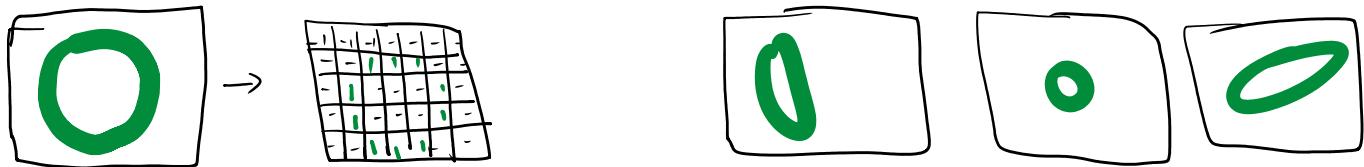
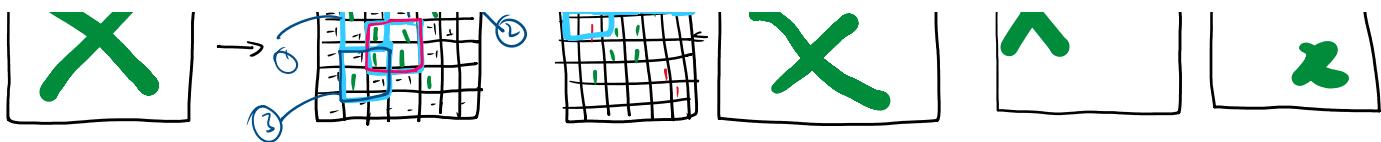


- ▷ CNN -
 - local area Network \rightarrow computationally less expensive
 - automatically learns spatial features in data

How CNN works

- - Convolutional layers
 - Relu layers
 - Pooling
 - F CNN -



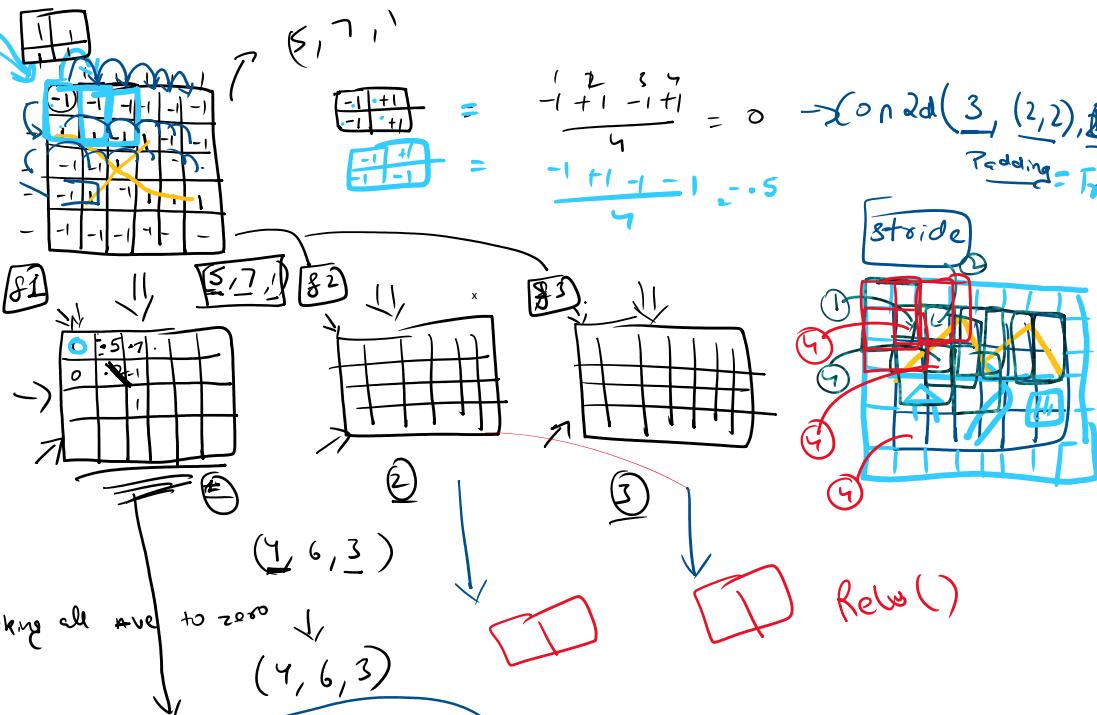


Convolution

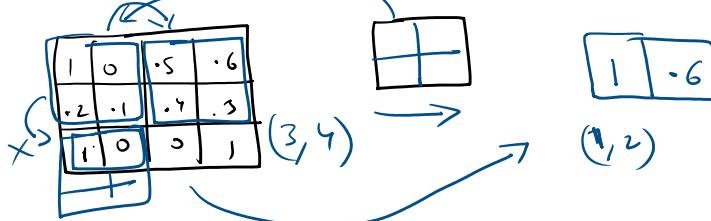
feature \rightarrow

n - size of image $(2, 2)$
 f - size of kernel $(2, 2)$
 $(n-f+1)$
 $(5-2+1) (7-2+1)$
 $(4, 6, 1)$
Convolution layer

$\sum_{i=1}^4 \sum_{j=1}^4 (-1)^{i+j} f_{ij} = \frac{-1+1-1+1}{4} = 0 \rightarrow \text{conv}(3, (2, 2), 1)$
 $\sum_{i=1}^4 \sum_{j=1}^4 (-1)^{i+j} f_{ij} = \frac{-1+1-1-1}{4} = -0.5$
 $\text{Padding} = P_{\text{out}}$

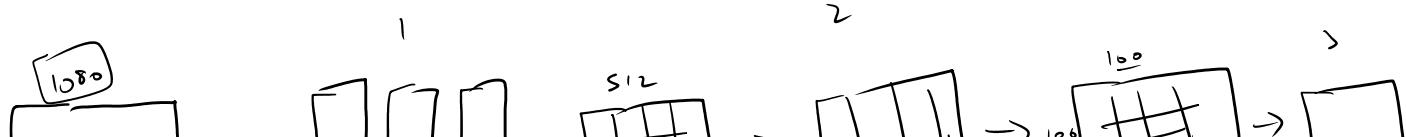
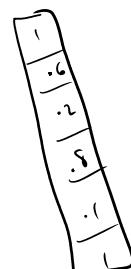


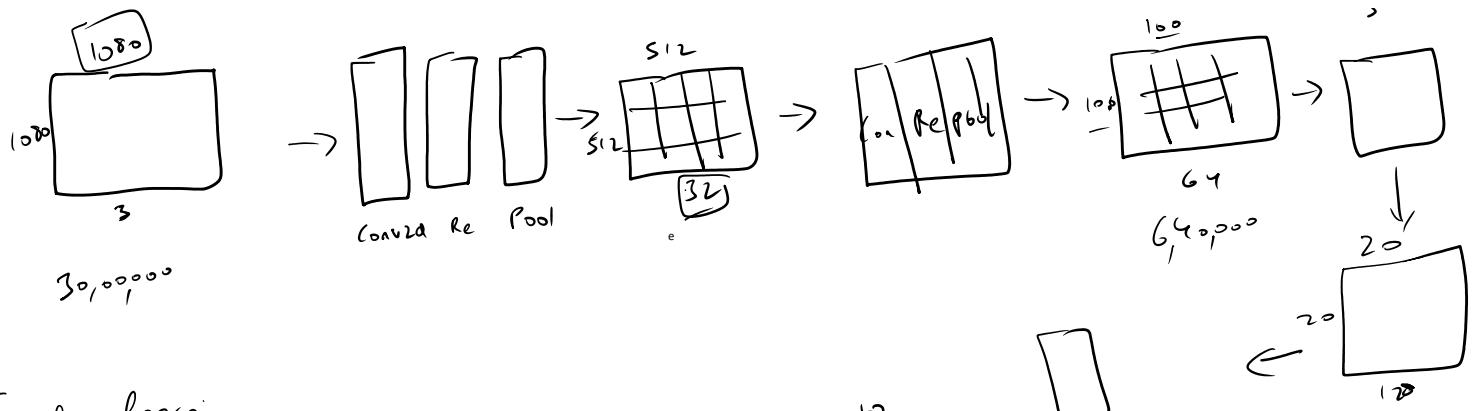
Max Pooling
Pooling layer :
 (Dimension red.)



Fraction

.1 .6
 .2 .8
 .1 .1





Transfer learning:

- Fine tuning :
 $\begin{cases} \text{Full} \\ \text{Partial} \end{cases}$
- Feature Extractor :

when to use
~~transfer learning~~

→ Why we need transfer learning :-?

- Reduce training time
- Avoid overfitting
- Data req. are less
- Performance

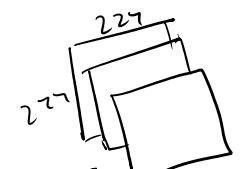
→ CNN-based Architectures :-

Application :- $(32, 32, 1)$

• LeNet-5 -
(1998) 2 Conv layers ($5,5$), tanh
 2 Avg. pooling

↓
2 Fully CNN with softmax

• AlexNet (2012) - ImageNet Challenge :-

5 conv. layer - relu $[224, 224, 3]$
3 maxpooling -
3 FCNN -
2 dropout -


» VGGNet (2014) - 19 layers

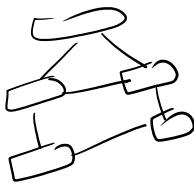
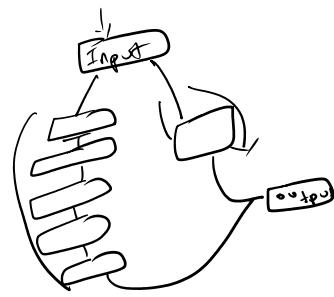
o GoogleNet (Inception)

o ResNet
2015

o DenseNet
2016

o MobileNet
2017

o EfficientNet



RCNN :- Region based CNN

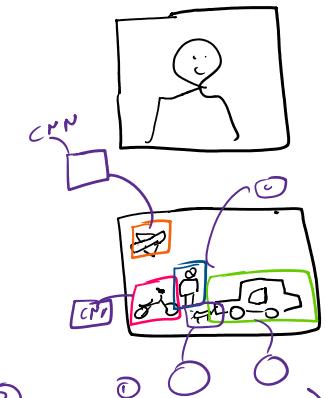
CNN - image clas

4 layers

→ [input shape = $(32, 32, 3)$]

1024 x
256
28

→ Region Proposal :-
→ Classification



Step 1:

- Selective Search :-

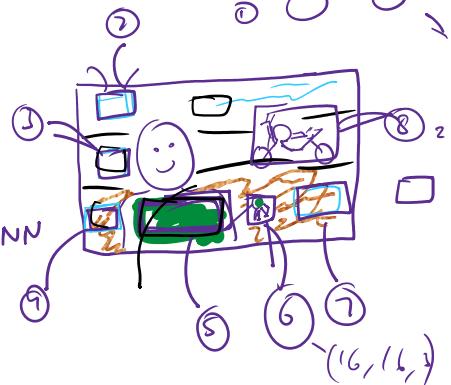
→ [2000]
2000

- color
- texture
- size
- shape

Step 2:

- Cropping / Reshaping to desired input of CNN

Step 3:-
Transfer Learning - Feature Extractor
- Classification



Step 4:- Bounding Box Regression :-

Disadvantages of RCNN

- Slow
- Training time
- Complex pipeline

- Fast RCNN $\rightarrow 24\uparrow$
- Faster RCNN $\rightarrow 20\uparrow$
- Masked RCNN
- Selective Search :- [2000]
 - Superpixel Segmentation

○ Superpixel Segmentation

Math intuition :-

$$\text{Similarity weight } \omega_{ij} = \frac{1}{|I(i) - I(j)|}$$

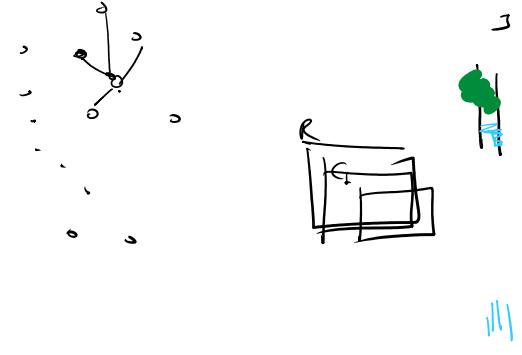
→ Color Histogram

→ Texture Histogram

• Size

• fill

2000

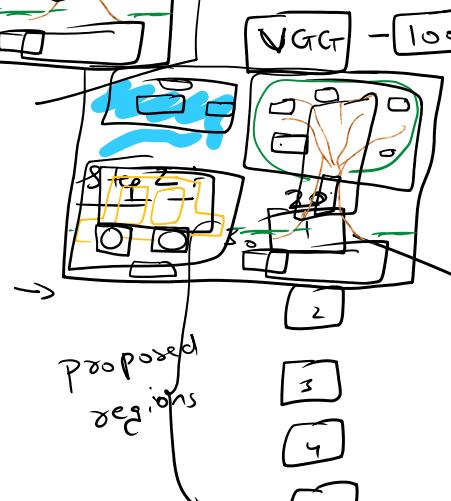
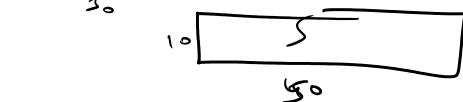
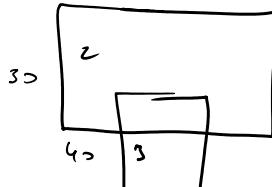
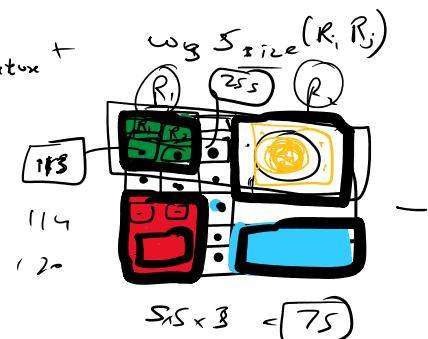
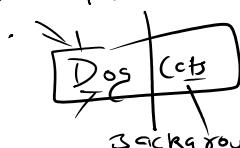


Similarity Score :-

$$\rightarrow S(R_i R_j) = \omega_1 S_{\text{color}}(R_i R_j) + \omega_2 S_{\text{texture}}(R_i R_j) + \omega_3 S_{\text{size}}(R_i R_j)$$

s = similarity score

R = pixel value



Step 3

→ ReNet

→ Verif

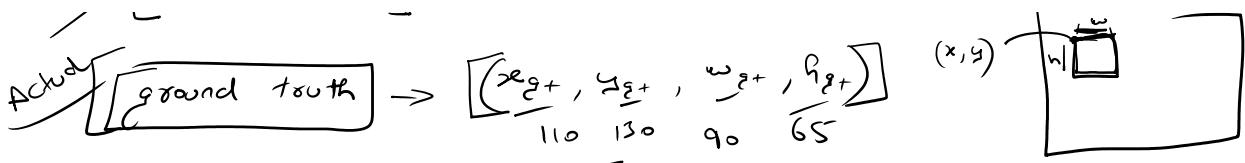
Alexnet

→ Classify

Step 4 :- Bounding Box Regression :-

$$\xrightarrow{\text{pred.}} (x, y, w, h) - \boxed{\text{pred.}} \quad (100, 20, 80, 60)$$

$$\xrightarrow{\text{actual ground truth}} [(x_g^+, y_g^+, w_g^+, h_g^+)] \quad (x, y)$$

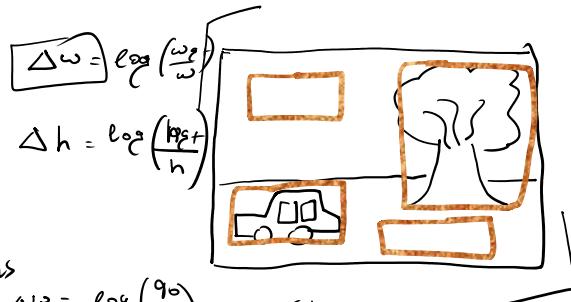


$$\Delta x = \frac{x_{g+} - x}{w}$$

$$\Delta y = \frac{y_{g+} - y}{h}$$

$$\Delta x = \frac{110 - 100}{80} = 0.125 \quad \Delta w = \log\left(\frac{w_g}{w}\right)$$

$$\Delta y = \frac{130 - 120}{60} = 0.167 \quad \Delta h = \log\left(\frac{h_g}{h}\right) = 0.034$$



$$x_{new} = x + \Delta x \cdot w \\ 100 + 0.125 \cdot 80$$

$$y_{new} = y + \Delta y \cdot h \\ w = w \cdot e^{\Delta w}$$

RCNN

- Region Propose
- Feature extract
- Classification
- Bounding Box Regress.

Fast RCNN

- o Single CNN
- o Region of Interest Pooling
Selective Search / R
- o Classification
- o Bounding Box Regression

Faster R-CNN

RPN

1 - CNN - $\boxed{1000}$

2 - feature map \rightarrow RPN

sliding to get region
objectiveness score

9000

4500

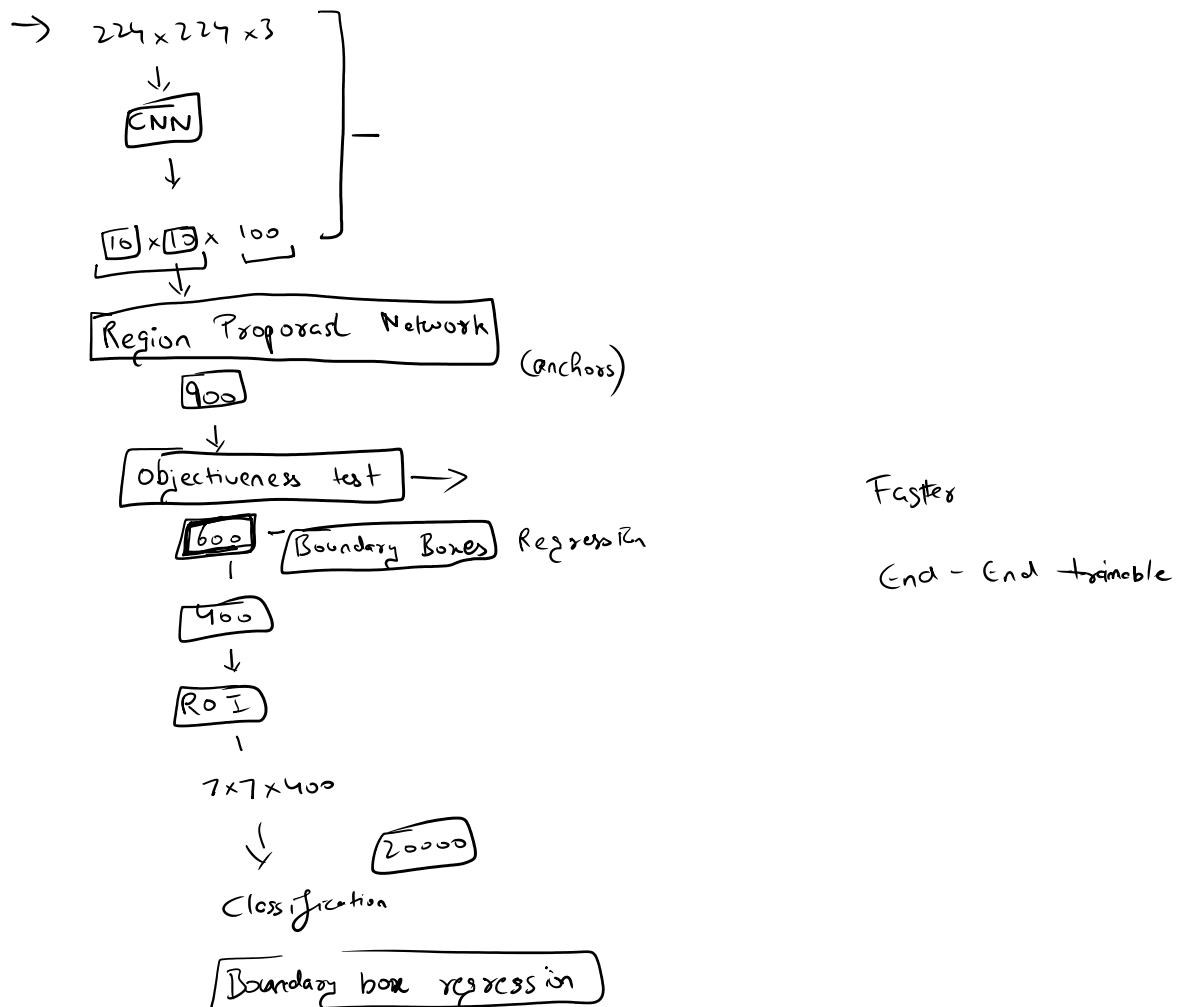
Bounding box

Non-maximum Suppression
 (7×7)

3) ROI Pooling

4) Classification

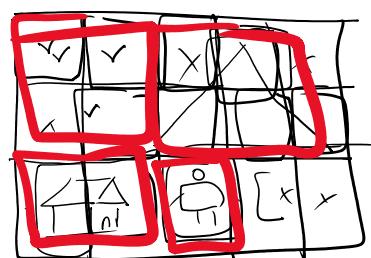
5) Bounding box regression



Yolo :- You Only Look Once

How Yolo works

- Create grid of $s \times s$ ($7 \times 7, 13 \times 3$)
- Bounding box & Confidence Score
- Class probabilities
- NMS - Non maximum suppression
- Speed & Accuracy Trade off



Limitation:

Variants:

◦ Yolov2: 2016

- Yolov2:
 - (Yolo 9009) • batch normalization
 - 416x416
 - 9000

- Yolov3:

◦ Yolov4:

- CSP-Darknet 53

◦ Yolo v5

◦ Ultralytics

- Yolov6

◦ deployment

◦ Yolov7 - 2022

- Yolov8 - 2023

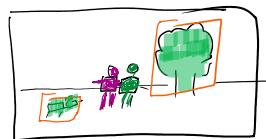
→ Object detection

→ Image segmentation

→ Object Tracking

Masked R-CNN :- Instance Segmentation

◦ boundary box - pixel level mask



Object Detection

◦ Faster R-CNN

- CNN
- RPN
- ROI Pooling
- FCNN/Regression

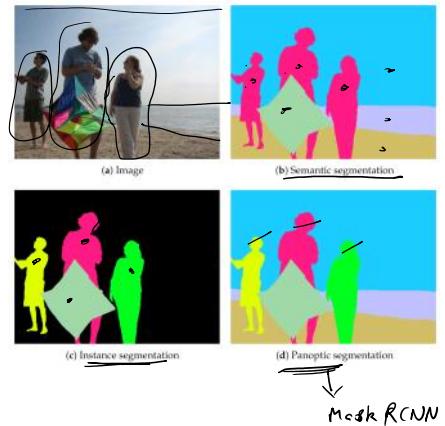
◦ Key Concepts

- Object detection
- Segmentation
- Instance Segmentation

- Segmentation
 - Instance Segmentation
 - ROI Pooling
 - FCNN / Regression

How Does it work?

- Feature Extraction ResNet
 - RPN - Regional Proposal Network
 - ROI Align
 - Classification / Bounding box Prediction
 - Mask Prediction



Use Cases !

- Autonomous Vehicles
 - Medical Imaging

Advantages

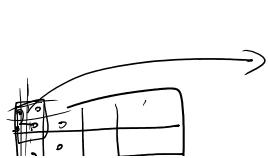
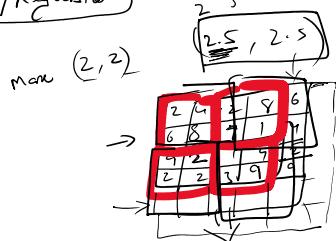
Faster RCNN

- ## • Feature Extraction (CNN)

◦ RPN

- ROI Pooling - Region of Interest

- FCNN / Regression



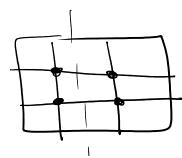
Masked R-CNN

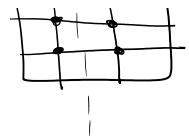
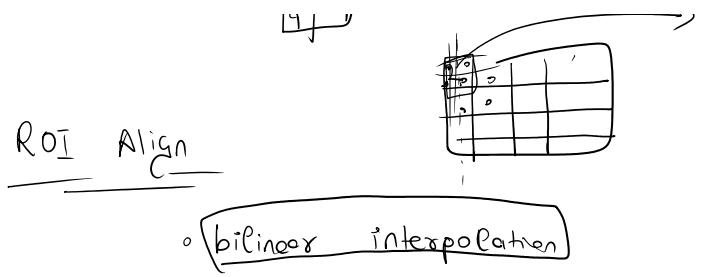
- ## ② Feature Extraction (CNN)

- o RPN

- ## • ROI Align

- Classification / Regression / Marked Head





$$v_{(x,y)} = (1-\alpha)(1-\beta)v_{(x_0,y_0)} + \alpha(1-\beta)v_{(x_1,y_0)} + (1-\alpha)\beta v_{(x_0,y_1)} + \alpha\beta v_{(x_1,y_1)}$$

Fine tuning :-

Yolo / Faster RCNN

- Setup Environment



Annotated dataset [Major time consumer]

- Prepare configuration file
- Start training

- Evaluate

Mixed RCNN

- Setup Environment

- Annotated dataset

- Registering dataset

- Start train

Detection

Object Tracking :-

- Single object tracking : (SOT)
- Multiple object tracking : (MOT)

→ Basic Components

- Detection
- Feature extraction
- Association
- Motion Prediction

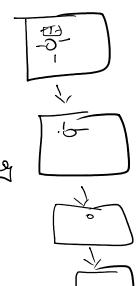
- Common techniques -

- Correlation-based tracking

- Optical Flow

- Kalman Filter

- Deep Learning based tracking



→ Challenges of Object tracking

- Occlusion :-

Applications :-



Occlusion :-

- Illumination Change ↗
- Scale Variation ↗
- Deformation ↗
- Background Clutter ↗

Applications :-

- Surveillance
- Robotics
- Object tracking
- Etc

• Working of OT ↗

◦ Detection :- CNN
Faster RCNN
Masked RCNN

◦ Motion Modelling ◦ Linear modelling $x_{t+1} = x_t + v_t \cdot \Delta t$

◦ Kalman Filter ↗

◦ predict-update cycle

→ Prediction Step ↗

◦ State transition ↗

$$x_t = F \cdot x_{t-1} + u_{t-1}$$

F = transition matrix

P = covariance matrix for uncertainty

Q = process noise

◦ Covariance Updater ↗

$$P_t = F \cdot P_{t-1} \cdot F^T + Q$$

H = measurement matrix

R = measurement noise

K_t = Kalman Gain

→ Update Step ↗

◦ Kalman Gain ↗

$$K_t = P_t \cdot H^T \cdot (H \cdot P_t \cdot H^T + R)^{-1}$$

◦ State Updater ↗

$$x_t = x_t + K_t \cdot (z_t - H \cdot x_t)$$

z_t = actual measurement at time t

Optical Flow :-

◦ pixel wise motion detection ↗

→ brightness constancy eqn ↗

$$I(x, y, t) = I(x + \Delta x, y + \Delta y, t + \Delta t)$$

Particle Filter (sequential Monte Carlo method)

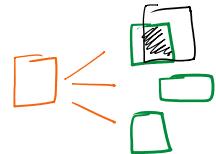
improved version of Kalman filter, can work on non-linear & non-gaussian processes

- Prediction
- Update
- Resampling = $w_t^i = p(z_t | x_t^i)$

Deep learning Modelling Methods

- SORT (Simple Online & Realtime Tracking)

bounding box detection + kalman filter



Stepst 1. Kalman filter predicts its next location

2. Detection in new frame & compared to predictions

$$\text{(Intersection over union)} = \text{IOU} = \frac{\text{Area of overlap}}{\text{Area of Union}}$$

3. Hungarian Algorithm minimize the total cost

- Deep SORT :

(Appearance based features)

$$\text{◦ Mahalanobis distance} = d(x,y) = \sqrt{(x-y)^T S^{-1} (x-y)}$$

Challenge Handling:

→ Occlusion Handling:

- Appearance based feature & Particle Filter

→ Data Association

Hungarian Algorithm & GRU (Gated Recurrent Unit)

→ 3D Object tracking

Stereo cameras, LiDAR

Additions

- Localization :

- Motion :

- o Motion Vectors

- o Tracking features :-
 - o SIFT (Scale invariant Feature Transform)
 - o Harris Corner Detector
 - o ORB (Oriented FAST & Rotated BRIEF)
- o Byte Tracking :-